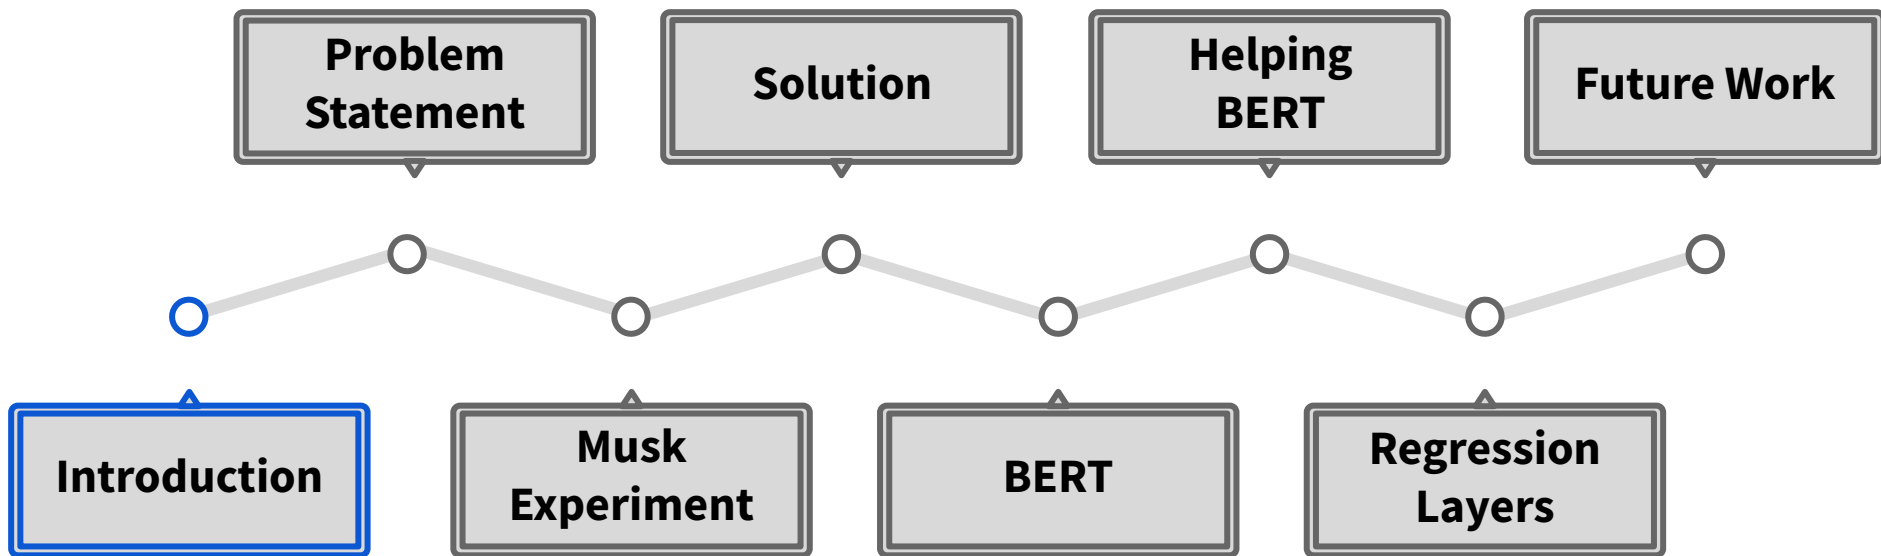


From a “lambo” to a Price Increase: Defining a Twitter Score for Financial Sentiment

Andrea Cicchini, Alexandru Ionut Pascariu, Victor Plesco



Why cryptocurrencies?




- Big Fluctuations in Price:
 - 1/25 - \$0.0083
 - 4/19 - \$0.4073 (+4807%)
- High Volatility due to Lack of Institutional Guarantor
 - “wild west”!
- Efficient Market Hypothesis (EMH)
 - Inefficient (N. A. Kyriazis, 2019)
- Social Media as Primary Source
 - Emotional Intelligence

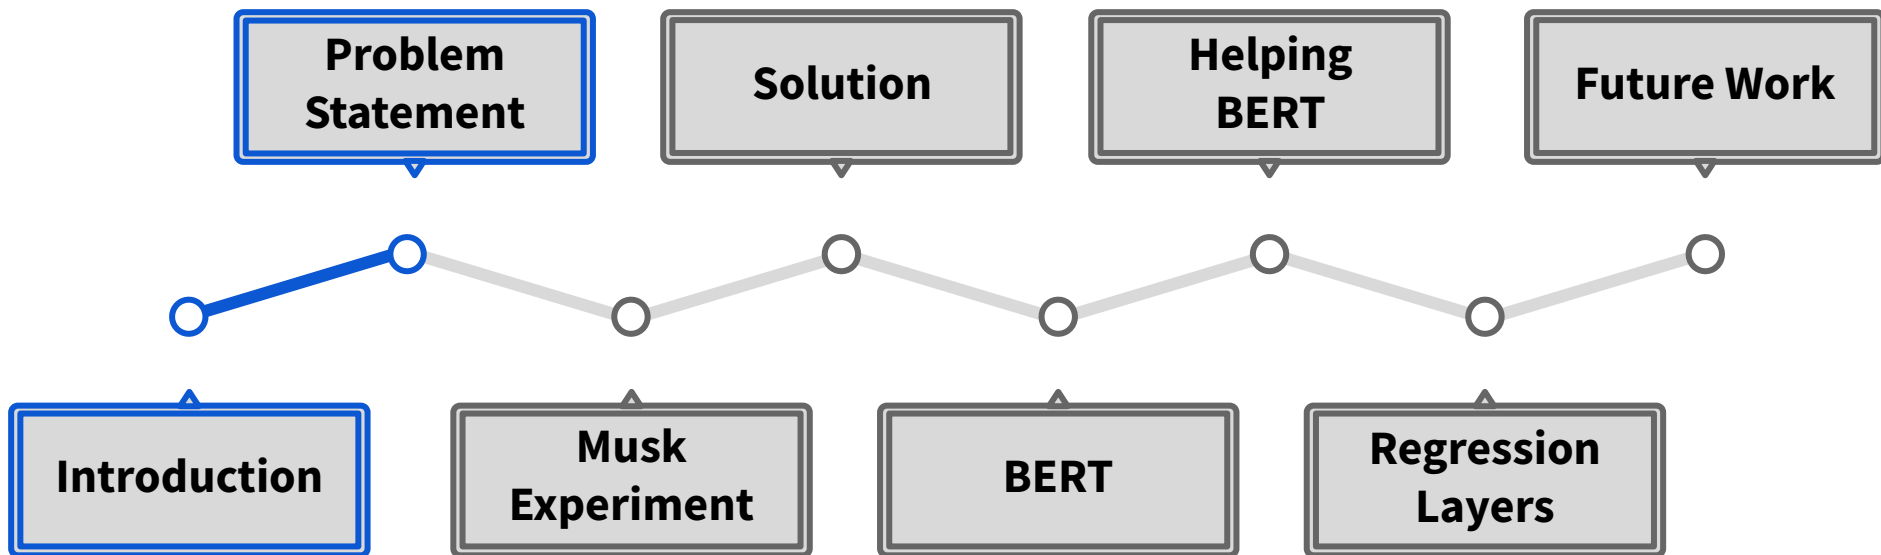


Social Network Sentiment For Financial Prediction

**Time Series
Analysis**



Sentiment Score



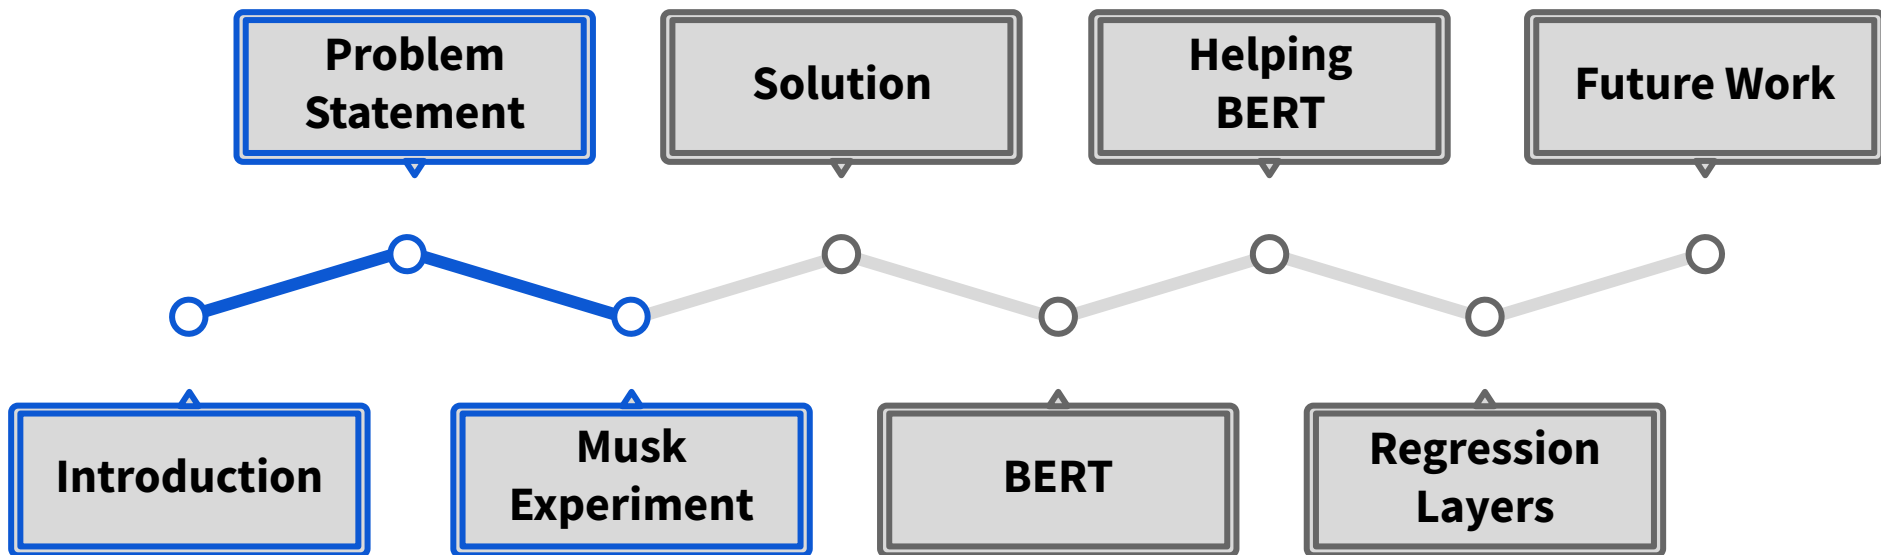
Is sentiment objective?

Which came first, the chicken or the egg?



or






Elon Musk, The Father of Dogecoin




Tweeting about Dogecoin (4th February)




Can Sentiment Analysis Detect It?


VADER: Valence Aware Dictionary and sEntiment Reasoner

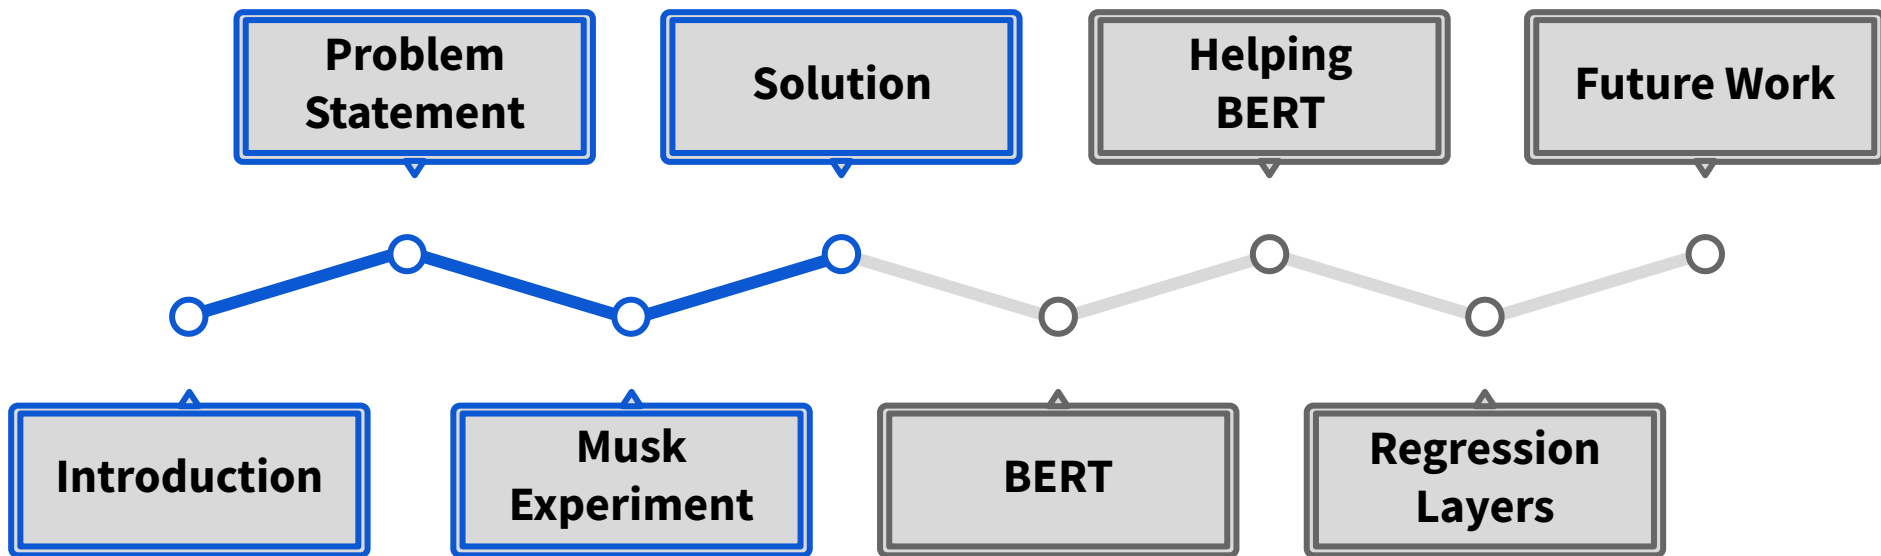


BERTweet: A pre-trained language model for English Tweets




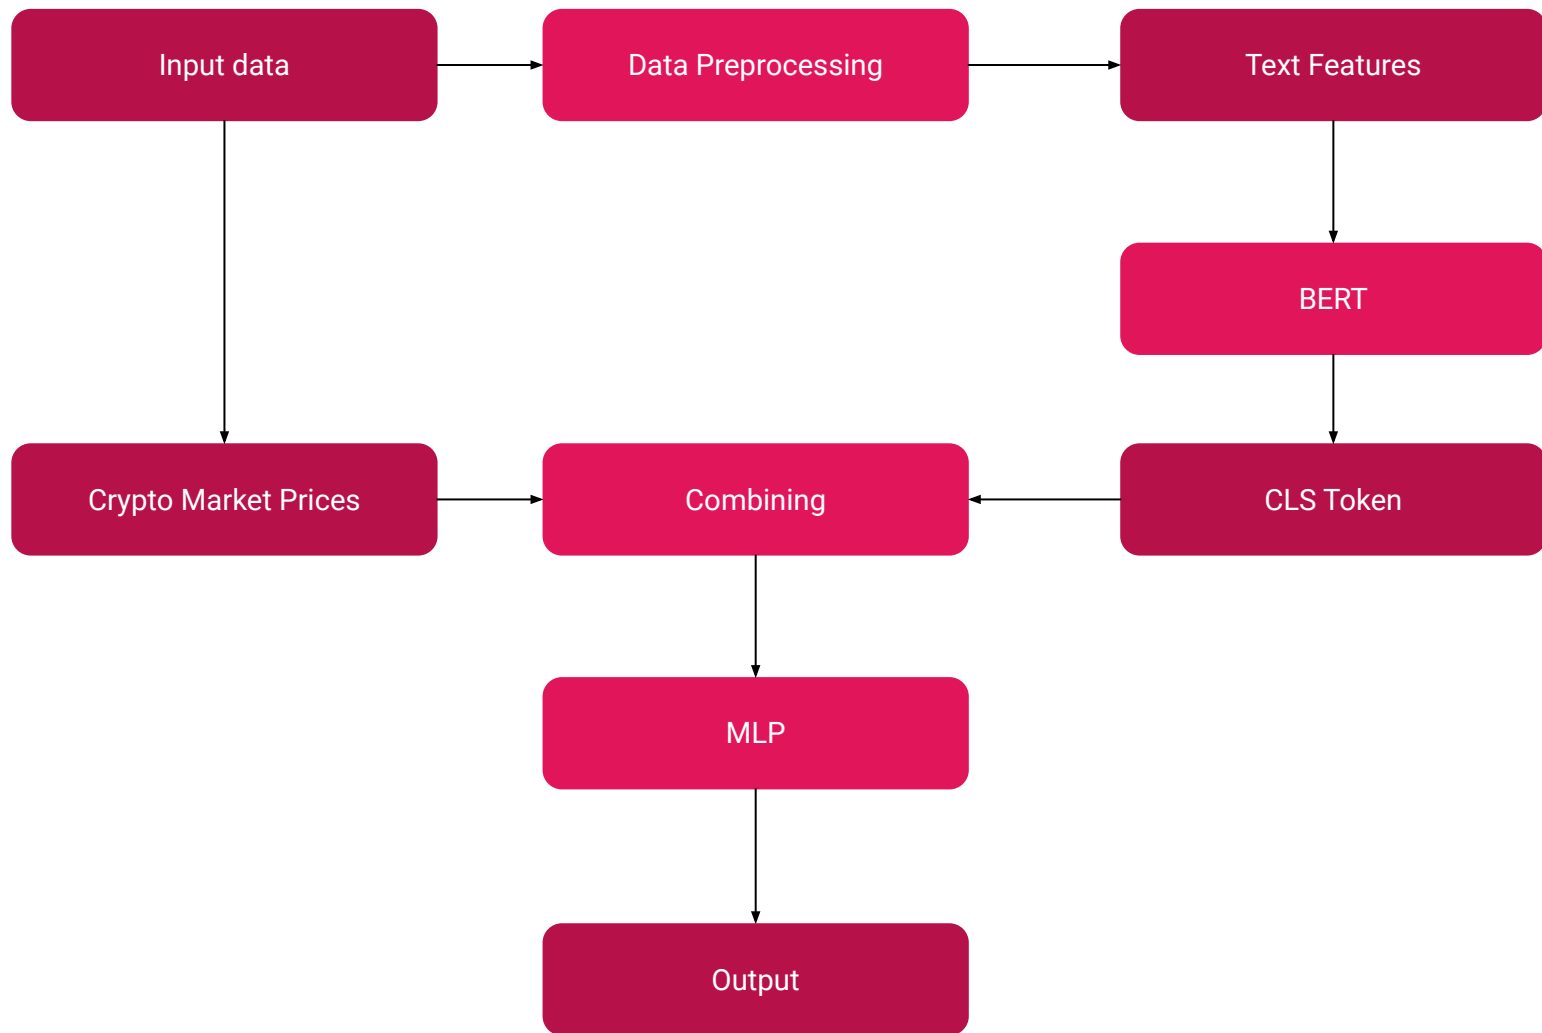
FinBERT: Financial Sentiment Analysis with BERT

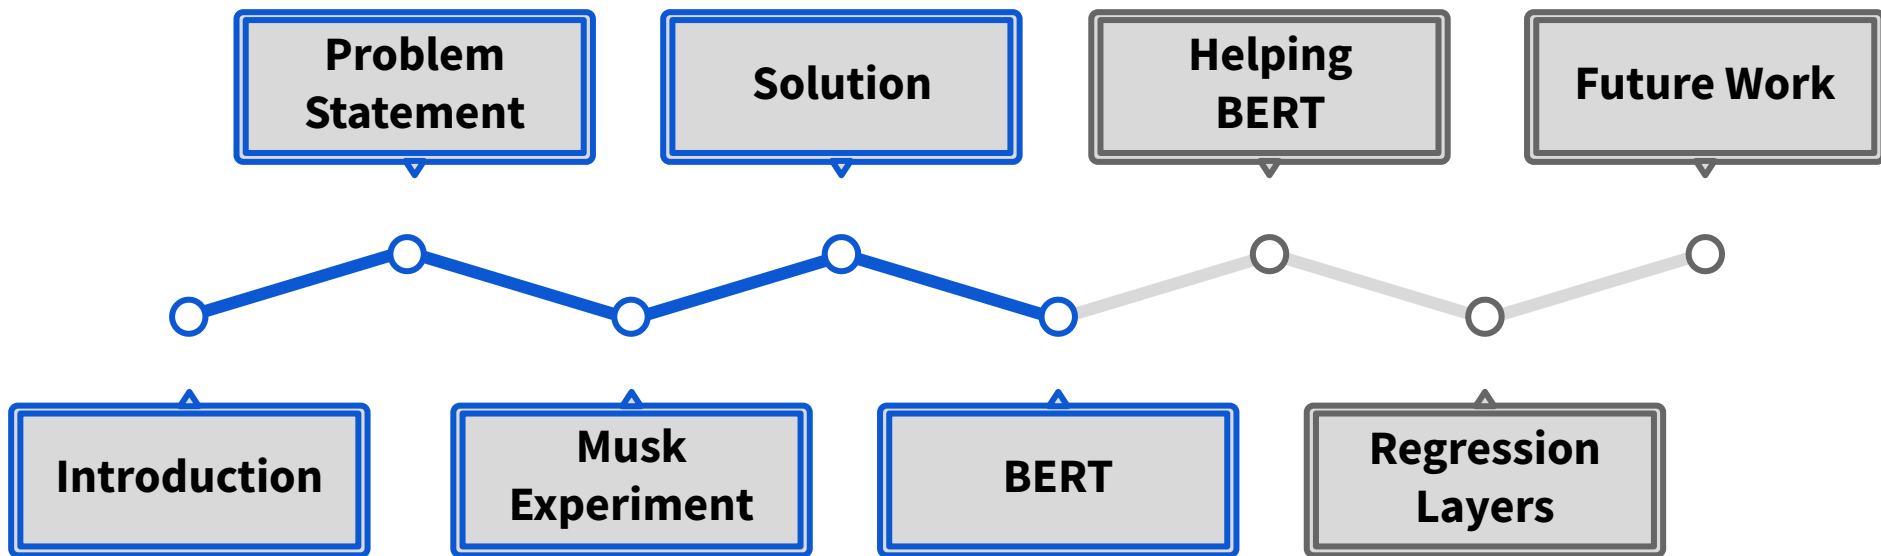





A New Sentiment Score: Twitter Score







BERT



- Bidirectional
- Encoder
- Representations
- Transformers

BERT Corpus & Parameters


For the pre-training corpus it was used:

→ The BooksCorpus (800M words) (Zhu et al., 2015)

→ English Wikipedia (2,500M words). It was ignored list, tables and header

Layer (type:depth-idx)	Param #
-----	-----
└BertEmbeddings: 1-1	--
└Embedding: 2-1	23,440,896
└Embedding: 2-2	393,216
└Embedding: 2-3	1,536
└LayerNorm: 2-4	1,536
└Dropout: 2-5	--
└BertEncoder: 1-2	--
└ModuleList: 2-6	--
└BertLayer: 3-1	7,087,872
└BertLayer: 3-2	7,087,872
└BertLayer: 3-3	7,087,872
└BertLayer: 3-4	7,087,872
└BertLayer: 3-5	7,087,872
└BertLayer: 3-6	7,087,872
└BertLayer: 3-7	7,087,872
└BertLayer: 3-8	7,087,872
└BertLayer: 3-9	7,087,872
└BertLayer: 3-10	7,087,872
└BertLayer: 3-11	7,087,872
└BertLayer: 3-12	7,087,872
└BertPooler: 1-3	--
└Linear: 2-7	590,592
└Tanh: 2-8	--
-----	-----
Total params:	109,482,240


BERT - Word Embedding and Tokens




BERT's Vocabulary

1. BERT is pre-trained → Vocabulary is fixed

2. Break down **unknown words** into **subwords**:




3. A subword exists for every word:




BERT - Word Embedding and Tokens

Types of Subword

All subwords start with “##” ...



Except for the first subword in a word.



Special Tokens

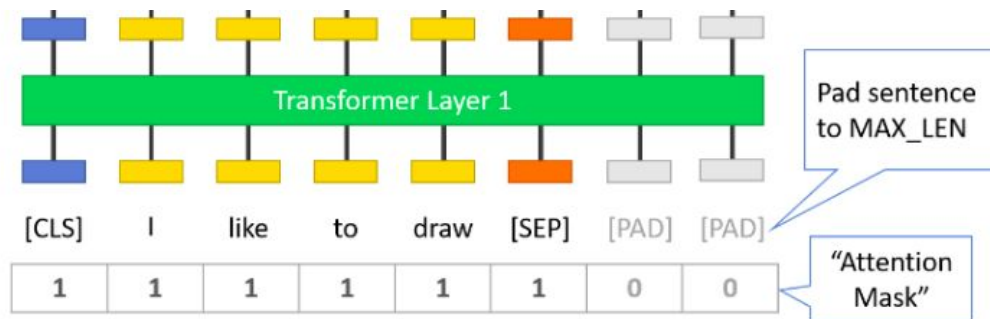
CLS → A special token representing the class of the input

SEP → A special token separating two different sentences in the same input

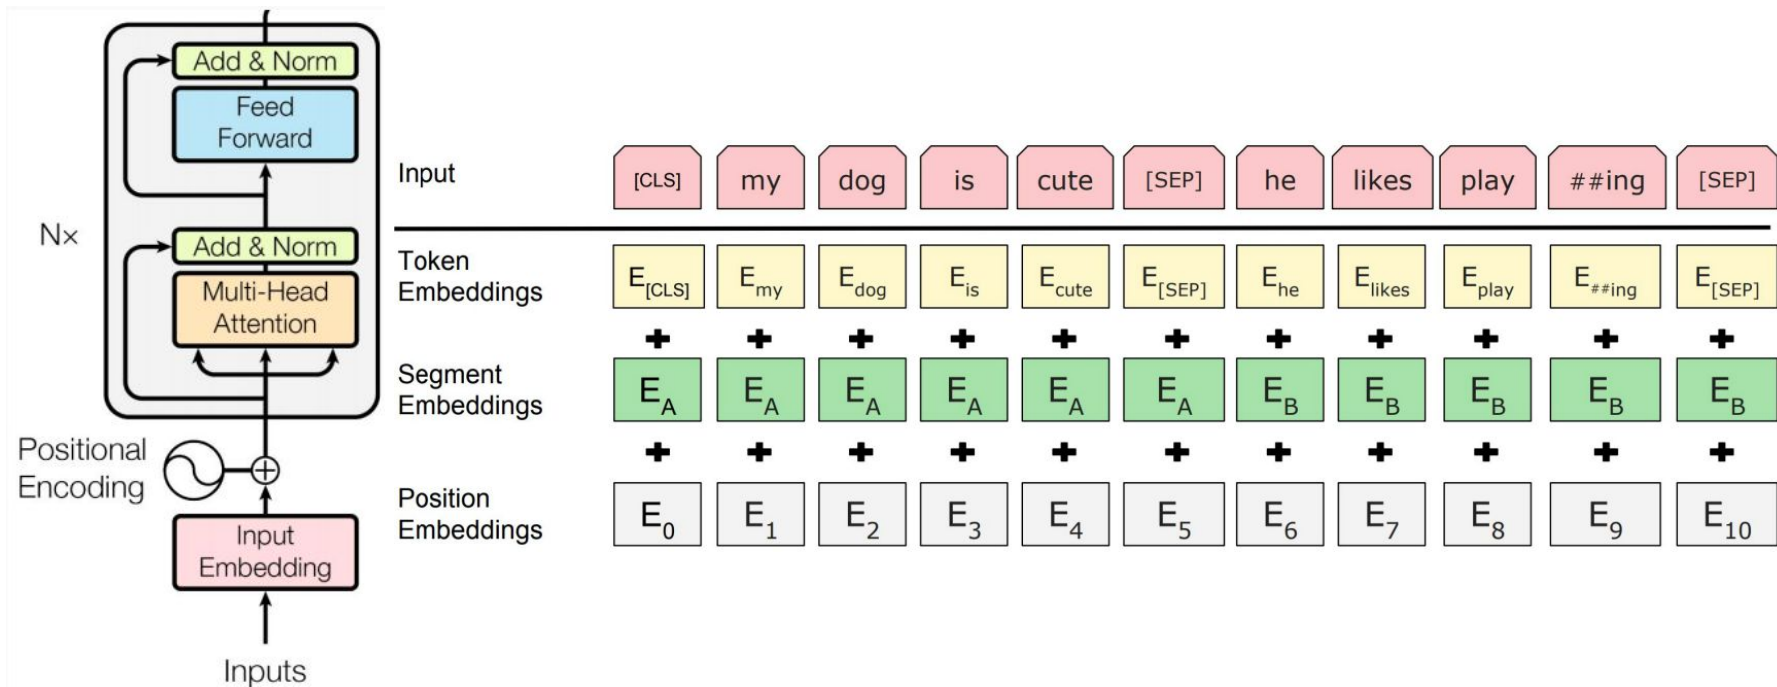
UNKNOWN → A special token representing an out-of-vocabulary token

PAD → A special token used to make arrays of tokens the same size for batching purpose. Will then be ignored by attention mechanisms or loss computation.


MASK → A special token representing a masked token



BERT - Inner Working



Position Embeddings




Linear

Fraction

$$\frac{1}{N-1}$$

Sentence 1





Sentence 2




Position Embeddings

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{\frac{2i}{d}}}\right)$$

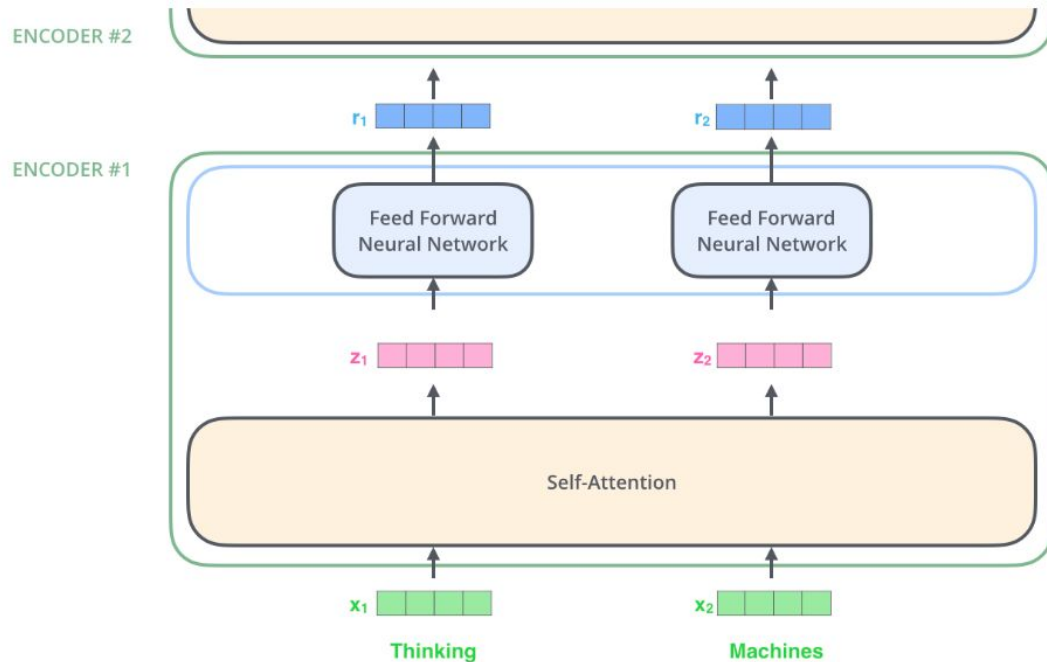
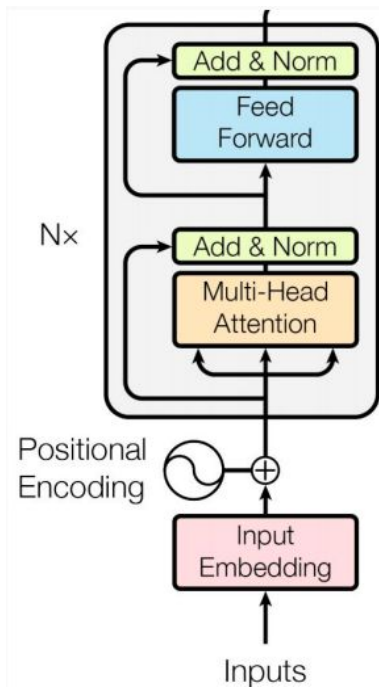
$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{\frac{2i}{d}}}\right)$$




Position Embeddings



Bert - Inner Working

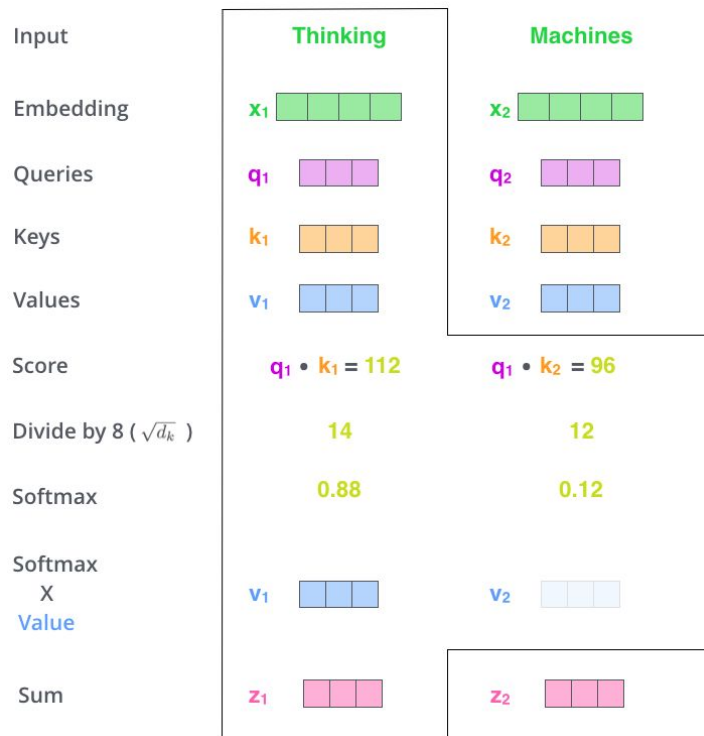


Bert - Inner Working Self Attention in Vector Form



Multiplying x_1 by the W^Q weight matrix produces q_1 , the "query" vector associated with that word. We end up creating a "query", a "key", and a "value" projection of each word in the input sentence.

Bert - Inner Working Self Attention in Vector Form



BERT - Inner Working Self Attention in Matrix Form

$$\begin{matrix} \text{X} \\ \begin{array}{|c|c|c|} \hline \square & \square & \square \\ \hline \square & \square & \square \\ \hline \end{array} \end{matrix} \times \begin{matrix} W^Q \\ \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \end{array} \end{matrix} = \begin{matrix} Q \\ \begin{array}{|c|c|} \hline \square & \square \\ \hline \square & \square \\ \hline \end{array} \end{matrix}$$

$$\begin{matrix} \text{X} \\ \begin{array}{|c|c|c|} \hline \square & \square & \square \\ \hline \square & \square & \square \\ \hline \end{array} \end{matrix} \times \begin{matrix} W^K \\ \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \end{array} \end{matrix} = \begin{matrix} K \\ \begin{array}{|c|c|} \hline \square & \square \\ \hline \square & \square \\ \hline \end{array} \end{matrix}$$

$$\begin{matrix} \text{X} \\ \begin{array}{|c|c|c|} \hline \square & \square & \square \\ \hline \square & \square & \square \\ \hline \end{array} \end{matrix} \times \begin{matrix} W^V \\ \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \end{array} \end{matrix} = \begin{matrix} V \\ \begin{array}{|c|c|} \hline \square & \square \\ \hline \square & \square \\ \hline \end{array} \end{matrix}$$

$$\text{softmax} \left(\frac{\begin{matrix} Q \\ \begin{array}{|c|c|} \hline \square & \square \\ \hline \square & \square \\ \hline \end{array} \end{matrix} \times \begin{matrix} K^T \\ \begin{array}{|c|c|} \hline \square & \square \\ \hline \square & \square \\ \hline \end{array} \end{matrix}}{\sqrt{d_k}} \right) \begin{matrix} V \\ \begin{array}{|c|c|} \hline \square & \square \\ \hline \square & \square \\ \hline \end{array} \end{matrix} = \begin{matrix} Z \\ \begin{array}{|c|c|c|} \hline \square & \square & \square \\ \hline \square & \square & \square \\ \hline \end{array} \end{matrix}$$


The self-attention calculation in matrix form

BERT - Inner Working Multi-Headed Attention

1) This is our
input sentence*

Thinking
Machines



2) We embed
each word*




* In all encoders other than #0,
we don't need embedding.
We start directly with the output
of the encoder right below this one



3) Split into 8 heads.
We multiply X or
 R with weight matrices



...



4) Calculate attention
using the resulting
 $Q/K/V$ matrices



...



5) Concatenate the resulting Z matrices,
then multiply with weight matrix W^O
to produce the output of the layer




...



W^O



Z

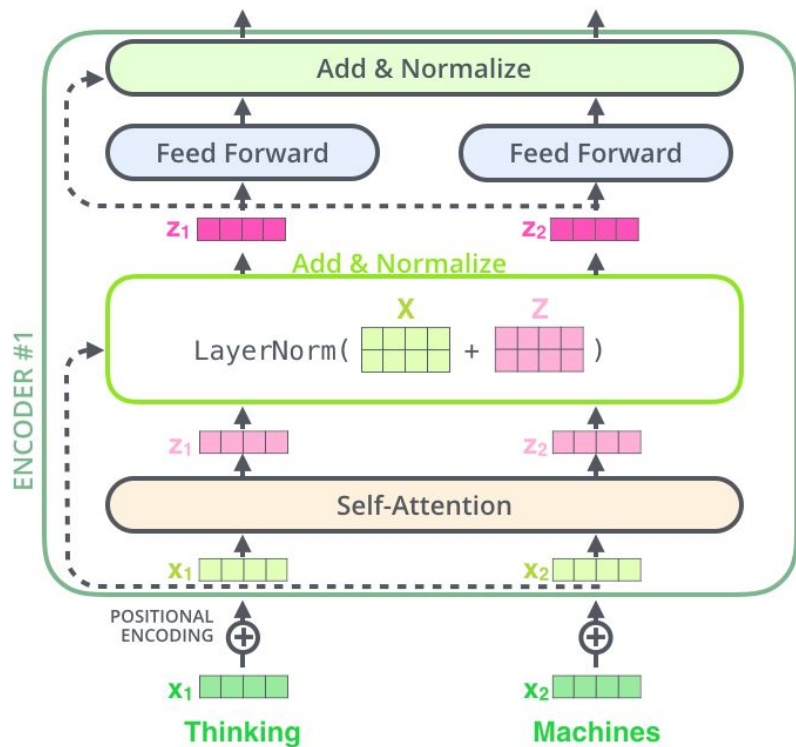


```

(bert): BertModel(
  (embeddings): BertEmbeddings(
    (word_embeddings): Embedding(30522, 768, padding_idx=0)
    (position_embeddings): Embedding(512, 768)
    (token_type_embeddings): Embedding(2, 768)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
  (encoder): BertEncoder(
    (layer): ModuleList(
      (0): BertLayer(
        (attention): BertAttention(
          (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)
            (value): Linear(in_features=768, out_features=768, bias=True)
            (dropout): Dropout(p=0.1, inplace=False)
          )
          (output): BertSelfOutput(
            (dense): Linear(in_features=768, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
          )
        )
      )
      (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
      )
      (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
    )
  )
)

```

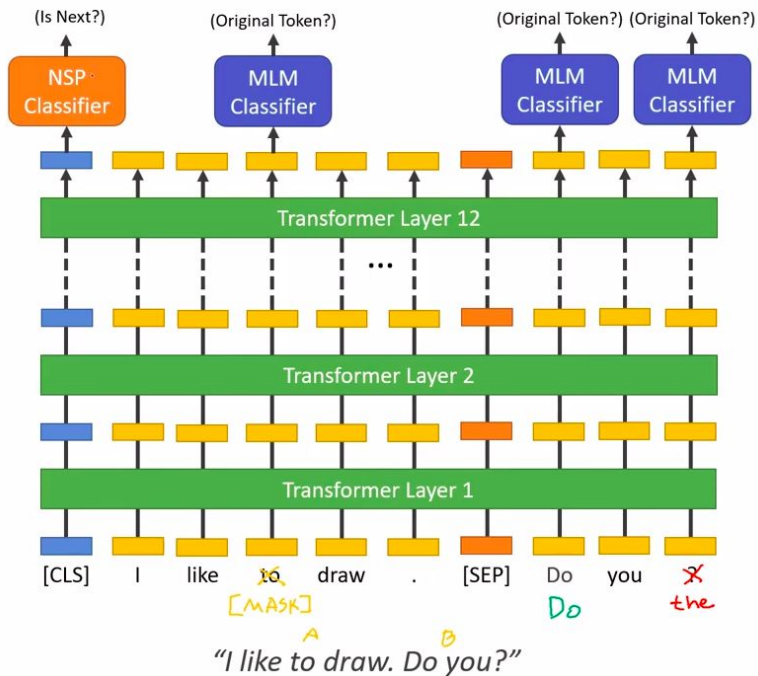

BERT - Add & Normalize




						mean (μ)	std (σ)
x_0 = When	0.98	1.28	0.41	0.27	0.41	0.67	0.44
x_1 = you	0.52	0.01	2.06	0.27	0.33	0.64	0.82
x_2 = play	2.22	0.27	0.10	0.41	2.06	1.01	1.04
x_3 = the	0.99	1.00	0.11	0.27	0.33	0.54	0.42
x_4 = game	0.52	0.01	0.33	2.06	0.52	0.69	0.79
x_5 = of	0.10	2.06	0.73	0.27	0.41	0.71	0.79
x_6 = thrones	0.33	0.01	0.13	0.27	1.28	0.40	0.51
	f_0	f_1	f_2	f_3	f_4		

$$x_i = \frac{x_i^d - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

BERT - Pre-Training Tasks



■ Untouched ■ Masked ■ Swapped ■ Swapped-Same

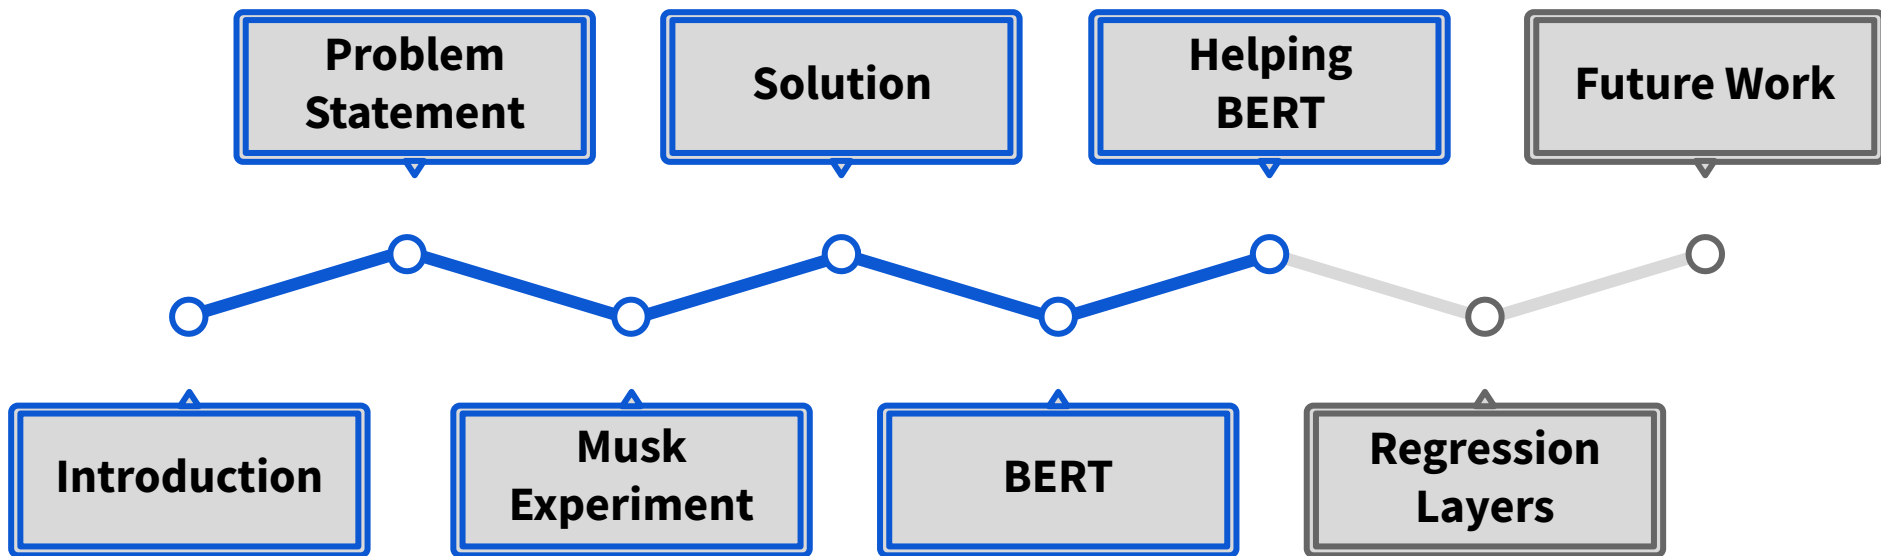


BERT - Shortcomings

BERT is very large → Embedding Layer with 24M weights (30K tokens each of 768 values)
→ Transformers (12) with 85M weights
→ slow fine-tuning
→ slow inferencing

Jargon → domain-specific language

Model		Parameters	Layers	Hidden	Embedding
BERT	base	108M	12	768	768
	large	334M	24	1024	1024
	xlarge	1270M	24	2048	2048



01

Help BERT with Informal Language

- Preprocessing Data

02

Slow Inferencing

- Dynamic Padding

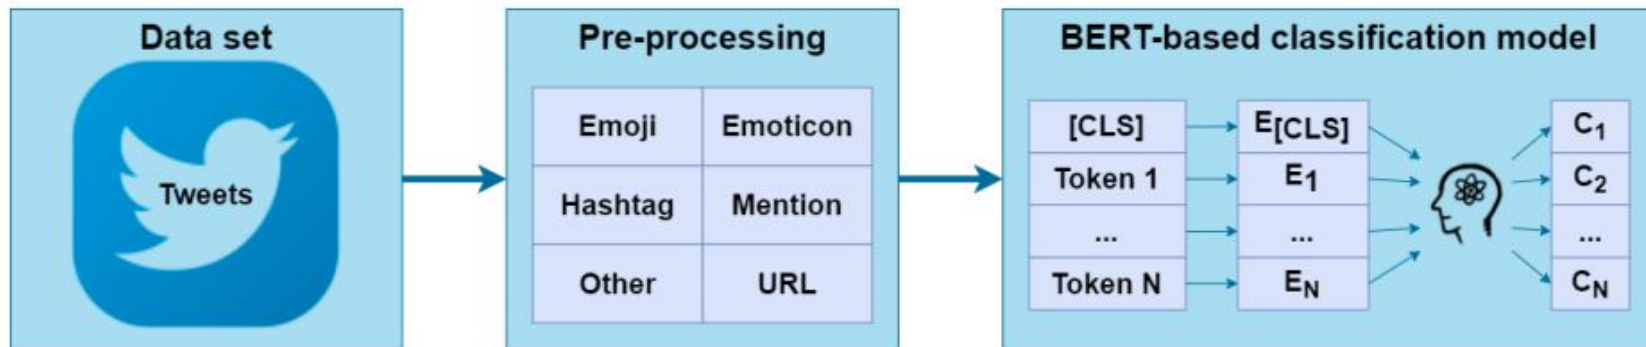
03

Heavy Data Management

- Batch Division

Data Preprocessing → Tweet to BERT

Tweets are very noisy and obscure particular slang people maintain on the creative use of social media. Tweets are characterized by a particular slang people maintain on the creative use of social media.



Data Preprocessing → Uninterpretable Words

According to the literature [4] we tried to preprocess the dataset with two particular aims:

REDUCE NOISE

- dates
- email
- money
- percentage
- url
- time
- phone
- number

THE “SPARKLY” TWITTER CREATIVITY

- user
- cashtag
- hashtag
- emoji

Data Preprocessing → Examples

- 30/02/2003 → <date>
- \$42 → <money>
- 120% → <percentage>
- 4:20PM → <time>
- 351.8744170 → <phone>
- 42 → <number>
- <https://da/S4m45aRu13z\n\nTrial> → <url>

Data Preprocessing → Examples


- @mybeautifulaccount → @user
- \$BTC \$BTCpizzamandolino → <cashtag> <cashtag>
- #lamboformambo → <lambo for mambo>
- 🤪 → Rolling on the Floor Laughing

Data preprocessing → Results

"- 💰 Buy now/pay later 🚀 #BlackFriday\n- 🌐 World's largest art lending platform coming to @Algorand.\nDigital asset exchanges coming to \$algo\n- #Algorand is interoperable 😄 \n- Faster than \$eth, better tech, green 🌱\nStart watching@ 5:42:00 ➡ <https://t.co/mtSxwVMUCT>\n#AlgoAutumn"

"- **money bag** Buy now/pay later **rocket** **<Black Friday>** - **globe** showing Americas World's largest art lending platform coming to . - Digital asset exchanges coming to **<cashtag>** - is interoperable star-struck - Faster than **<cashtag>** better tech, green seedling Start watching@ **<time>** right arrow **<url>**"

Preprocessing Arrays



Fixed Padding Length

	1	2	3	4	5	6	7	8	9	10	11	12	13	14
1	_Eh	_bien	_c	'	_est	_un	_bon	indicateu	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]
2	_Ouais	_je	_suis	_un	_coureur	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]
3	_Ils	_ne	_sont	_pas	important	_	.	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]
4	_Il	_y	_a	_de	nombreux	condition	_qui	_ne	_sont	_pas	_visibles	_	.	[PAD]
5	_Chaque	_zone	_de	_l	'	_île	_offre	_quelque	_chose	_de	_différent	_	.	[PAD]
6	_Mais	_tu	_peux	_vivre	_avec	_eux	_	.	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]
7	_Un	_grand	_homme	_	,	_dit	-	il	_	.	[PAD]	[PAD]	[PAD]	[PAD]
8	_Elle	_a	_été	_menée	_en	_silence	_	.	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]
9	_Tu	er	beaucoup	_de	_fourmis	_de	_feu	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]
10	_La	_question	_est	_de	_savoir	_si	_clin	ton	_a	_le	_cul	ot	_	.
11	_C	'	_est	_vrai	_	.	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]
12	_Dans	_ce	_domaine	_	,	_seuls	_les	_sa	ther	i	_le	_savent	_	.

Batch Length: 14

Batch Length: 14

Batch Length: 14

Total Tokens: 168

Dynamic Padding


	1	2	3	4	5	6	7	8	9	10	11	12	13	14	
1	_Eh	_bien	_c	'	_est	_un	_bon	indicateur	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]		Batch Length: 13
2	_Ouais	_je	_suis	_un	_coureur	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]		
3	_Ils	_ne	_sont	_pas	important	_	.	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]		
4	_Il	_y	_a	_de	nombreux	condition	_qui	_ne	_sont	_pas	_visibles	_	.		
5	_Chaque	_zone	_de	_l	'	_île	_offre	_quelque	_chose	_de	_différent	_	.		Batch Length: 13
6	_Mais	_tu	_peux	_vivre	_avec	_eux	_	.	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]		
7	_Un	_grand	_homme	_	,	_dit	-	il	_	.	[PAD]	[PAD]	[PAD]		
8	_Elle	_a	_été	_menée	_en	_silence	_	.	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]		
9	_Tu	er	beaucoup	_de	_fourmis	_de	_feu	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	Batch Length: 14
10	_La	_question	_est	_de	_savoir	_si	_clin	ton	_a	_le	_cul	ot	_	.	
11	_C	'	_est	_vrai	_	.	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]	
12	_Dans	_ce	_domaine	_	,	_seuls	_les	_sa	ther	i	_le	_savent	_	.	


Total Tokens: 160

Uniform Length Batching (Our Approach)

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	
2	_Ouais	_je	_suis	_un	_coureur	[PAD]	[PAD]								Batch Length: 7
11	_C	'	_est	_vrai	_	.	[PAD]								
3	_Ils	_ne	_sont	_pas	important	_	.								
9	_Tu	er	beaucoup	_de	_fourmis	_de	_feu								
1	_Eh	_bien	_c	'	_est	_un	_bon	indicateur	[PAD]	[PAD]					Batch Length: 10
6	_Mais	_tu	_peux	_vivre	_avec	_eux	_	.	[PAD]	[PAD]					
8	_Elle	_a	_été	_menée	_en	_silence	_	.	[PAD]	[PAD]					
7	_Un	_grand	_homme	_	,	_dit	-	il	_	.					
5	_Chaque	_zone	_de	_l	'	_île	_offre	_quelque	_chose	_de	_différent	_	.	[PAD]	Batch Length: 14
4	_Il	_y	_a	_de	nombreux	condition	_qui	_ne	_sont	_pas	_visibles	_	.	[PAD]	
10	_La	_question	_est	_de	_savoir	_si	_clin	ton	_a	_le	_cul	ot	_	.	
12	_Dans	_ce	_domaine	_	,	_seuls	_les	_sa	ther	i	_le	_savent	_	.	

Total Tokens: **124**






Impact of PAD Token on Accuracy

Padding Strategy	Model	Batch Size	Max Len	Test Accur.	GPU	Training Time per Epoch (mm:ss)
Smart Batching	BERT-base	16	400	0.935	Tesla K80	0:35:06
Fixed Padding	BERT-base	16	400	0.93	Tesla K80	0:53:14

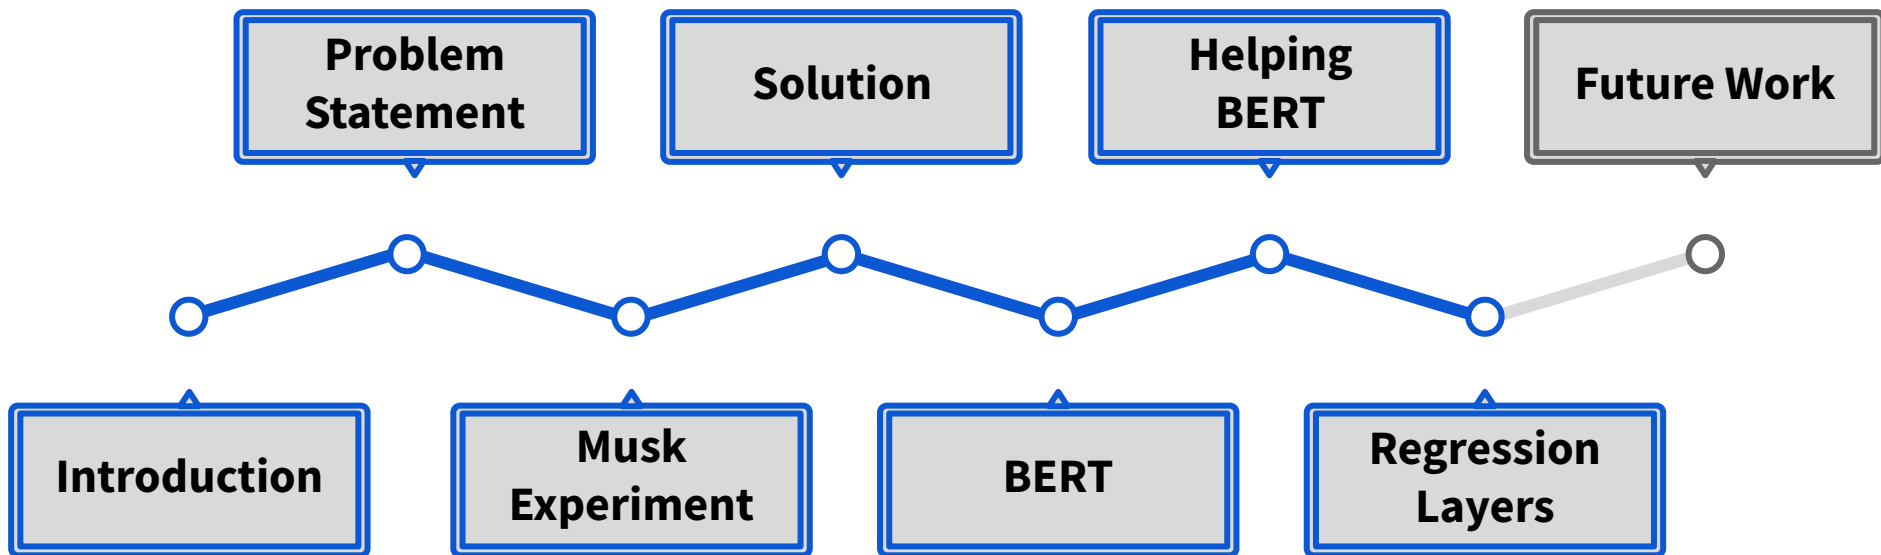
Our Data in Numbers

Coin	Tweets
ADA	59918
AVAX	18049
DOGE	35976
DOT	22158
ETH	291569
LTC	11856
SOL	183511
UNI	5695
XRP	20651

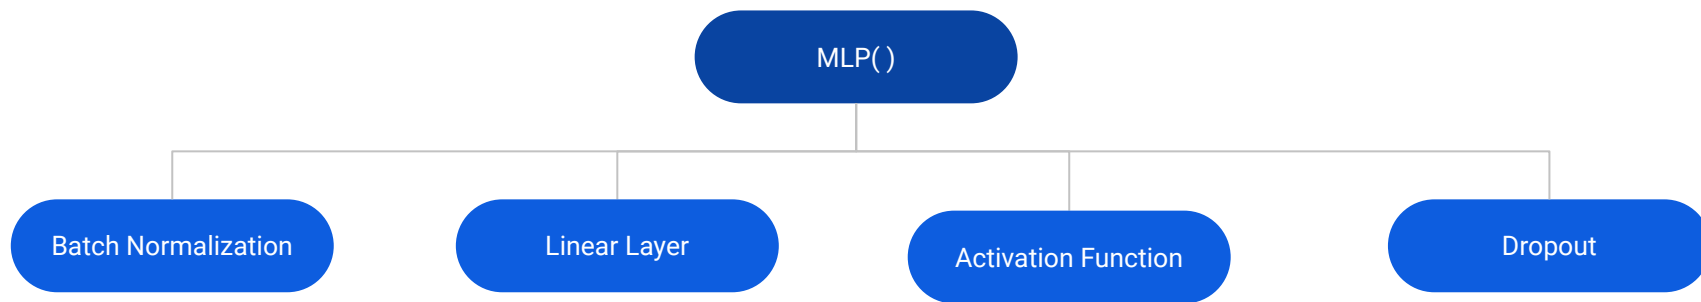
Train, Test, Validation Split



Data Batch Division



Multi Layer Perceptron



First Attempt

```

hyperparams = pd.DataFrame(data={"n_in_1": [768, 768, 768],
                                "n_out_1": [512, 151, 768],
                                "batchnorm_1": [False, False, False],
                                "activ_1": [nn.ReLU, nn.ReLU, nn.ReLU],
                                "dropout_p_1": [None, None, None],

                                "n_in_2": [512, 151, 768],
                                "n_out_2": [341, 341, 512],
                                "batchnorm_2": [True, True, True],
                                "activ_2": [nn.LeakyReLU, nn.LeakyReLU, nn.LeakyReLU],
                                "dropout_p_2": [0.4, 0.5, 0.6],

                                "n_in_3": [341, 341, 512],
                                "n_out_3": [227, 512, 341],
                                "batchnorm_3": [True, True, True],
                                "activ_3": [nn.LeakyReLU, nn.LeakyReLU, nn.LeakyReLU],
                                "dropout_p_3": [0.4, 0.5, 0.6],

                                "n_in_4": [227, 512, 341],
                                "n_out_4": [151, 768, 227],
                                "batchnorm_4": [True, True, True],
                                "activ_4": [nn.LeakyReLU, nn.LeakyReLU, nn.LeakyReLU],
                                "dropout_p_4": [0.4, 0.5, 0.6],

                                "n_in_5": [151, 768, 227],
                                "n_out_5": [1, 1, 1],
                                "batchnorm_5": [True, True, True],
                                "activ_5": [nn.LeakyReLU, nn.LeakyReLU, nn.LeakyReLU],
                                "dropout_p_5": [None, None, None],


                                "learning_rate": [0.0001, 0.001, 0.01]})

```

hyperparams


Training Parameters

Optimizer	→	ADAM
Loss Function	→	MSE
Batch Size	→	64
Early Stopping	→	patience=4




Second Attempt

Michael Nielsen's Tip




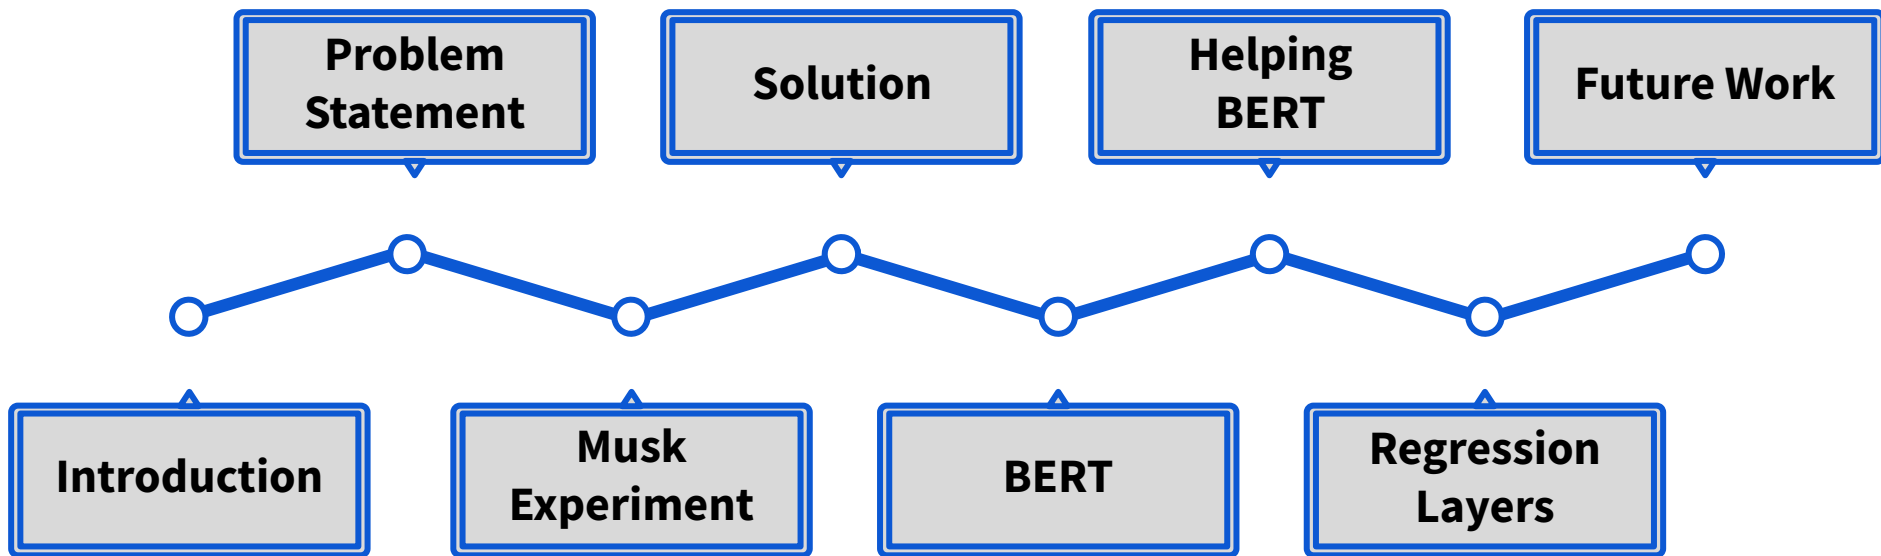
1st Quartile



4th Quartile







Future Work: No Results is a Result :)

- No Correlation Discovered:
 - Propose a different preprocessing or work on a pre-trained model on Tweets.
- Handle Twitter Bots (~80%):
 - -- Krypto-Invaders - -\n\n🎮 - from 0.002 \$ETH\n\n🧩 - only 300. Ever!! - 1:1 NFTs\n\n👤 - Playable Game
\n\nCollect here:\n\n🌐👉 <https://t.co/HcGFC6euWB>\n\n\n#pixelnft #galxy #NFTcollectibles #nftgaming #nftgame <https://t.co/LyvZQ79HgE>
- Try on a different lag

References

- [1] Jacob Devlin Ming-Wei Chang Kenton Lee Kristina Toutanova Google AI Language. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding
- [2] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Lukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. In Advances in Neural Information Processing Systems, pages 6000–6010.
- [3] Stephane Lathuilière, Pablo Mesejo, Xavier Alameda-Pineda, Member IEEE, and Radu Horaud. 2020. A Comprehensive Analysis of Deep Regression
- [4] Marco Pota, Mirko Ventura, Rosario Catelli and Massimo Esposi. 2021. An Effective BERT-Based Pipeline for Twitter Sentiment Analysis: A Case Study in Italian
- [5] Kostadin Yotov, Emil Hadzhikolev, Stanka Hadzhikoleva. 2020. Determining the Number of Neurons in Artificial Neural Networks for Approximation, Trained with Algorithms Using the Jacobi Matrix
- [6] STEFANO CRESCI, FABRIZIO LILLO, DANIELE REGOLI, SERENA TARDELLI and MAURIZIO TESCONI. 2018. Cashtag piggybacking: uncovering spam and bot activity in stock microblogs on Twitter
- [7] Chris McCormick, Smart Batching Tutorial - Speed Up BERT Training:
<https://mccormickml.com/2020/07/29/smart-batching-tutorial/#s5-fine-tune-bert>
- [8] Anthony Galtier. 2021. Fine-tuning BERT for a regression task: is a description enough to predict a property's list price?
- [9] Chris McCormick YouTube channel: <https://www.youtube.com/c/ChrisMcCormickAI/videos>
- [10] Hedü Math of Intelligence YouTube channel: <https://www.youtube.com/c/HeduMathematicsofIntelligence/videos>
- [11] Neural Network and Deep Learning, free online book: <http://neuralnetworksanddeeplearning.com/index.html>